



# Kubernetes Installation and Deployment



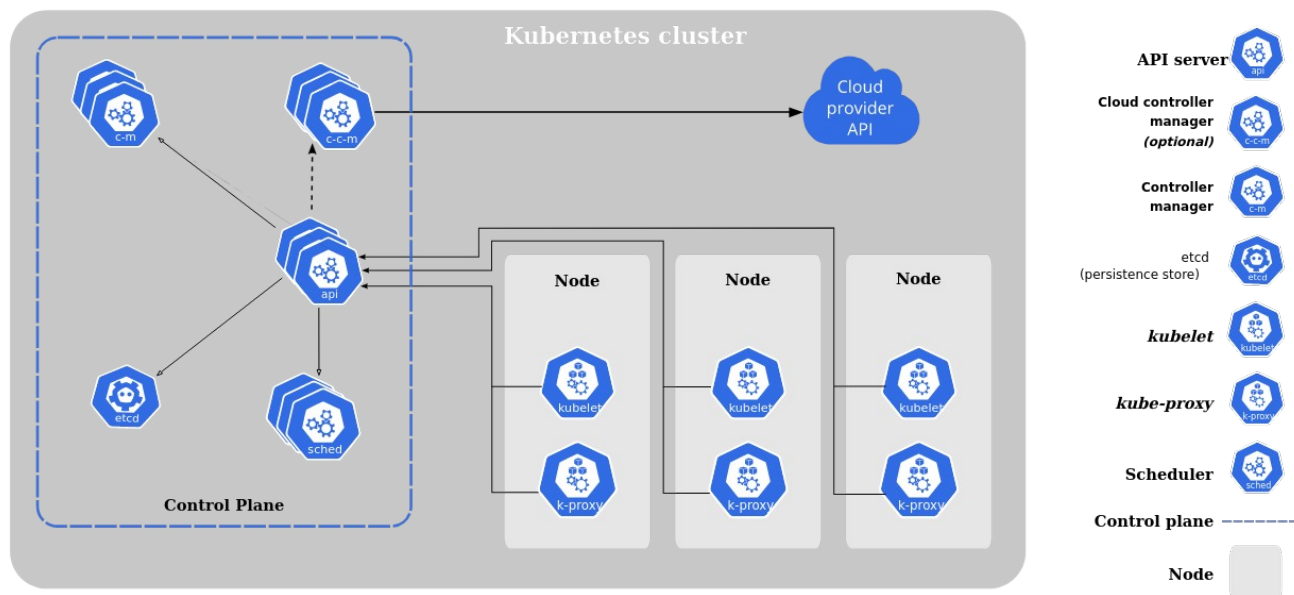
## Objectives

Chapter 1	Introduction
Chapter 2	Installation Prerequisite
Chapter 3	Installation on Centos 7
Chapter 4	Creating the Cluster
Chapter 5	Creating Docker Image
Chapter 6	Test Deployment
Chapter 7	Creating the Service
Chapter 8	Final Output
Chapter 9	References

**1.Introduction:** Kubernetes is a portable, extensible, open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation. It has a large, rapidly growing ecosystem. Kubernetes services, support, and tools are widely available.

**Kubernetes Cluster:** A Kubernetes cluster consists of a set of worker machines, called nodes, that run containerized applications. Every cluster has at least one worker node.

Basic architecture of a cluster with different components.



## 2.Installation Prerequisite:

- ✓ A compatible Linux host with 2 GB or more of RAM & 2 CPUs or more.
- ✓ Full network connectivity between all machines in the cluster (public or private network)
- ✓ Docker in the machine.

## Different Installation Toolbox:

- ◆ Kubectl → Bare Metal Installation.
- ◆ MiniKube → Virtualized Environment for Kubernetes
- ◆ Kops → Kubernetes on AWS
- ◆ Kubernetes on GCP → Kubernetes running on Google Cloud Platform
- ◆ AKS → Azure Kubernetes Services on Azure platform.

**Installation Packages:** Below packages need to install on the all machines.

- ◆ Kubectl → the command to bootstrap the cluster.
- ◆ Kubelet → the component that runs on all of the machines in the cluster and does things like starting pods and containers.
- ◆ Kubectl → the command line utility to talk to the cluster.

**Note:** kubectl will not install or manage kubelet or kubectl.



## Kubernetes Installation and Deployment



**3.Installation on Centos 7:** In this setup one virtual machine will prepare as master and another machine will play the role of slave. First install docker and Kubernetes on both machines.

#turn of the swap

```
$swapoff -a
```

# Comment out swap line in fstab so that it remains disabled after reboot

```
$vi /etc/fstab
```

#common utilities install

```
$yum install -y yum-utils device-mapper-persistent-data lvm2
```

#it is preferred to disable the firewall

```
$sudo systemctl disable --now firewalld
```

#if needs to active the firewall than add the following ports in the firewall on Master

```
$sudo firewall-cmd --add-port={6443,2379-2380,10250,10251,10252,5473,179,5473}/tcp --permanent
```

```
$sudo firewall-cmd --add-port={4789,8285,8472}/udp --permanent
```

```
$sudo firewall-cmd --reload
```

#for worker nodes

```
sudo firewall-cmd --add-port={10250,30000-32767,5473,179,5473}/tcp --permanent
```

```
sudo firewall-cmd --add-port={4789,8285,8472}/udp --permanent
```

```
sudo firewall-cmd --reload
```

#Set the net.bridge.bridge-nf-call-iptables to '1' in sysctl config file. This ensures that packets are properly processed by IP tables during filtering and port forwarding

```
$cat <<EOF > /etc/sysctl.d/k8s.conf
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
EOF
```

```
$sysctl --system
```

#Enable br\_netfilter Kernel Module

```
$sudo modprobe overlay
```

```
$sudo modprobe br_netfilter
```

```
$sudo tee /etc/sysctl.d/kubernetes.conf<<EOF
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOF
```

```
$sudo sysctl --system
```

#in this installation Docker run will used. So need to install the docker run time.

# Install packages

```
$sudo yum install -y yum-utils device-mapper-persistent-data lvm2
```

```
$sudo yum-config-manager --add-repo https://download.docker.com/linux/centos/docker-ce.repo
```



## Kubernetes Installation and Deployment



```
$sudo yum install docker-ce docker-ce-cli containerd.io
```

```
# Create required directories
```

```
$sudo mkdir /etc/docker
```

```
$sudo mkdir -p /etc/systemd/system/docker.service.d
```

```
# Create daemon json config file
```

```
$sudo tee /etc/docker/daemon.json <<EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2",
  "storage-opts": [
    "overlay2.override_kernel_check=true"
  ]
}
EOF
```

```
# Start and enable Services
```

```
$sudo systemctl daemon-reload
```

```
$sudo systemctl restart docker
```

```
$sudo systemctl enable docker
```

```
#add the repository for the Kubernetes.
```

```
$cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-`$basearch`
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
exclude=kubelet kubeadm kubectl
EOF
```

```
# Set SELinux in permissive mode (effectively disabling it)
```

```
$sudo setenforce 0
```

```
$sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

```
#install kubelet kubeadm kubectl
```

```
$sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
```

```
#check the version
```

```
$sudo kubeadm version
```



# Kubernetes Installation and Deployment



```
[root@worker1 ~]# kubeadm version
kubeadm version: &version.Info{Major:"1", Minor:"23", GitVersion:"v1.23.4", GitCommit:"e6c093d87ea4cbb530a7b2ae91e54c0842d8308a", Git
TreeState:"clean", BuildDate:"2022-02-16T12:36:57Z", GoVersion:"go1.17.7", Compiler:"gc", Platform:"linux/amd64"}
[root@worker1 ~]# docker version
Client: Docker Engine - Community
Version: 20.10.12
API version: 1.41
Go version: go1.16.12
Git commit: e91ed57
Built: Mon Dec 13 11:45:41 2021
OS/Arch: linux/amd64
Context: default
Experimental: true

Server: Docker Engine - Community
Engine:
Version: 20.10.12
API version: 1.41 (minimum version 1.12)
Go version: go1.16.12
Git commit: 459d0df
Built: Mon Dec 13 11:44:05 2021
OS/Arch: linux/amd64
Experimental: false
containerd:
Version: 1.4.12
GitCommit: 7b11cfaabd73bb80907dd23182b9347b4245eb5d
runc:
Version: 1.0.2
GitCommit: v1.0.2-0-g52b36a2
docker-init:
Version: 0.19.0
GitCommit: de40ad0
```

```
[root@master ~]# kubeadm version
kubeadm version: &version.Info{Major:"1", Minor:"23", GitVersion:"v1.23.4", GitCommit:"e6c093d87ea4cbb530a7b2ae91e54c0842d8308a", Git
TreeState:"clean", BuildDate:"2022-02-16T12:36:57Z", GoVersion:"go1.17.7", Compiler:"gc", Platform:"linux/amd64"}
[root@master ~]# docker version
Client: Docker Engine - Community
Version: 20.10.12
API version: 1.41
Go version: go1.16.12
Git commit: e91ed57
Built: Mon Dec 13 11:45:41 2021
OS/Arch: linux/amd64
Context: default
Experimental: true

Server: Docker Engine - Community
Engine:
Version: 20.10.12
API version: 1.41 (minimum version 1.12)
Go version: go1.16.12
Git commit: 459d0df
Built: Mon Dec 13 11:44:05 2021
OS/Arch: linux/amd64
Experimental: false
containerd:
Version: 1.4.12
GitCommit: 7b11cfaabd73bb80907dd23182b9347b4245eb5d
runc:
Version: 1.0.2
GitCommit: v1.0.2-0-g52b36a2
docker-init:
Version: 0.19.0
GitCommit: de40ad0
```

**Note:** After installation reboot the both machines

#start and enable the kubelet service

\$sudo systemctl start kubelet && sudo systemctl enable kubelet

#Change the cgroup-driver

\$sudo docker info | grep -i cgroup



## Kubernetes Installation and Deployment



#need to make sure the docker-ce and kubernetes are using same 'cgroup'

```
sed -i 's/cgroup-driver=systemd/cgroup-driver=cgroupfs/g' /etc/systemd/system/kubelet.service.d/10-kubeadm.conf
```

### 4.Creating the cluster:

#pull the container images so that it could easy crate the cluster during offline.

```
$sudo kubeadm config images pull
```

```
[root@master ~]# systemctl enable kubelet
[root@master ~]# sudo kubeadm config images pull
[config/images] Pulled k8s.gcr.io/kube-apiserver:v1.23.4
[config/images] Pulled k8s.gcr.io/kube-controller-manager:v1.23.4
[config/images] Pulled k8s.gcr.io/kube-scheduler:v1.23.4
[config/images] Pulled k8s.gcr.io/kube-proxy:v1.23.4
[config/images] Pulled k8s.gcr.io/pause:3.6
[config/images] Pulled k8s.gcr.io/etcd:3.5.1-0
[config/images] Pulled k8s.gcr.io/coredns/coredns:v1.8.6
```

#initialize the master

```
$sudo kubeadm init --pod-network-cidr=10.10.0.0/16 --upload-certs --control-plane-endpoint=192.168.0.131
```

```
Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

You can now join any number of the control-plane node running the following command on each as root:

kubeadm join 192.168.0.131:6443 --token hqdxay.uz1li02llutez7ki \
--discovery-token-ca-cert-hash sha256:9bb9e16fd6249affb8b5d29191aab473a94c2c7ef279091912c9a626cb458d43 \
--control-plane --certificate-key 85c1851b27333da668ff2b5415297a5a58ffee557ae31b221f31c99d799922fe

Please note that the certificate-key gives access to cluster sensitive data, keep it secret!
As a safeguard, uploaded-certs will be deleted in two hours; If necessary, you can use
"kubeadm init phase upload-certs --upload-certs" to reload certs afterward.

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 192.168.0.131:6443 --token hqdxay.uz1li02llutez7ki \
--discovery-token-ca-cert-hash sha256:9bb9e16fd6249affb8b5d29191aab473a94c2c7ef279091912c9a626cb458d43
```

#Configure kubectl using commands in the output

```
$mkdir -p $HOME/.kube
```

```
$sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
$sudo chown $(id -u):$(id -g) $HOME/.kube/config
```



# Kubernetes Installation and Deployment



#check the cluster status

\$kubectl cluster-info

```
[root@master ~]# kubectl cluster-info
Kubernetes control plane is running at https://192.168.0.131:6443
CoreDNS is running at https://192.168.0.131:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

#deploy the flannel network to the kubernetes cluster using the kubectl command

\$kubectl apply -f <https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml>

#joining the worker node with master

\$kubeadm join 192.168.0.131:6443 --token 4kwlkq.0qko5na74eb7y5r1 --discovery-token-ca-cert-hash sha256:225c197a3a592689aa4d4cbb2abfc9e1974c64194249137b16f3ef8e2417a737

```
[root@worker1 ~]# kubeadm join 192.168.0.131:6443 --token 4kwlkq.0qko5na74eb7y5r1 \
> --discovery-token-ca-cert-hash sha256:225c197a3a592689aa4d4cbb2abfc9e1974c64194249137b16f3ef8e2417a737
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

#to check all nodes status

\$kubectl get nodes

#to check all pods

\$kubectl get pods --all-namespaces

```
[root@master ~]# kubectl get nodes
NAME        STATUS    ROLES          AGE    VERSION
master      Ready     control-plane,master  24m    v1.23.4
worker1     Ready     <none>          19m    v1.23.4

[root@master ~]# kubectl get pods --all-namespaces
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  coredns-64897985d-gfm5j                1/1     Running   0           24m
kube-system  coredns-64897985d-jdlg5                1/1     Running   0           24m
kube-system  etcd-master                            1/1     Running   1           24m
kube-system  kube-apiserver-master                   1/1     Running   1           24m
kube-system  kube-controller-manager-master          1/1     Running   2           24m
kube-system  kube-flannel-ds-mdfsh                   1/1     Running   0           19m
kube-system  kube-flannel-ds-wpdkh                   1/1     Running   0           20m
kube-system  kube-proxy-jjqhz                        1/1     Running   0           24m
kube-system  kube-proxy-w4f8p                        1/1     Running   0           19m
kube-system  kube-scheduler-master                   1/1     Running   2           24m
tigera-operator  tigera-operator-59fc55759-lldjg        1/1     Running   0           22m
```





## Kubernetes Installation and Deployment



### 5. Creating Docker image:

#Dockerfile configuration file

```
FROM httpd:2.4
```

```
WORKDIR /usr/local/
```

```
EXPOSE 80
```

```
RUN sed -i "s/DirectoryIndex index.html/DirectoryIndex index.php/g" apache2/conf/httpd.conf
```

```
RUN apt-get update -y \
```

```
&& apt-get install git -y \
```

```
&& rm -rf apache2/htdocs/index.html \
```

```
&& git clone https://github.com/shindesharad71/Club-Manager.git apache2/htdocs/
```

#docker image build

```
$sudo docker build -t azizur013/club-manager .
```

#push the image into docker-hub

```
$sudo docker push azizur013/club-manager
```

```
Cloning into 'apache2/htdocs'...
Removing intermediate container ce263807d6b6
--> 5045a513d1f7
Successfully built 5045a513d1f7
Successfully tagged azizur013/club-manager:latest
[root@worker1 ~]# docker push azizur013/club-manager
Using default tag: latest
The push refers to repository [docker.io/azizur013/club-manager]
An image does not exist locally with the tag: azizur013/club-manager
[root@worker1 ~]# docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
azizur013/club-manager	latest	5045a513d1f7	53 seconds ago	245MB
httpd	2.4	faed93b28859	17 hours ago	144MB
k8s.gcr.io/kube-apiserver	v1.23.4	62930710c963	13 days ago	135MB
k8s.gcr.io/kube-proxy	v1.23.4	2114245ec4d6	13 days ago	112MB
k8s.gcr.io/kube-controller-manager	v1.23.4	25444908517a	13 days ago	125MB
k8s.gcr.io/kube-scheduler	v1.23.4	aceacb6244f9	13 days ago	53.5MB
rancher/mirrored-flannelcnf-flannel	v0.16.3	8cb5de74f107	4 weeks ago	59.7MB
rancher/mirrored-flannelcnf-flannel-cni-plugin	v1.0.1	ac40ce625740	5 weeks ago	8.1MB
k8s.gcr.io/etcd	3.5.1-0	25f8c7f3da61	3 months ago	293MB
k8s.gcr.io/coredns/coredns	v1.8.6	a4ca41631cc7	4 months ago	46.8MB
k8s.gcr.io/pause	3.6	6270bb605e12	6 months ago	683kB

```
[root@worker1 ~]# docker push azizur013/club-manager
Using default tag: latest
The push refers to repository [docker.io/azizur013/club-manager]
3d5ddba01f66: Pushed
1bcf3a329274: Pushed
0b9913462291: Pushed
bf25e956b68a: Pushed
1ecfa38fbb05: Pushed
db97effbf0f3: Pushed
1401df2b50d5: Pushed
latest: digest: sha256:55558865c7570d33ba0c178399526ec72fdaa0519e369d0f7413e0fbd588eda3 size: 1785
```





## Kubernetes Installation and Deployment



### 6.Test Deployment:

#manifest for the deployment

apiVersion: apps/v1

kind: Deployment

metadata:

name: test-deploy

labels:

environment: test

spec:

replicas: 3

selector:

matchLabels:

environment: test

minReadySeconds: 10

strategy:

rollingUpdate:

maxSurge: 1

maxUnavailable: 0

type: RollingUpdate

template:

metadata:

labels:

environment: test

spec:

containers:

- name: club-manager-application

image: azizur013/club-manager:latest

ports:

- containerPort: 80

#deployment command

\$kubectl apply -f club-manager-deployment.yml

```
[root@master ~]# kubectl apply -f club-manager-deployment.yml
deployment.apps/test-deploy created
[root@master ~]# kubectl get all
NAME                                     READY   STATUS    RESTARTS   AGE
pod/test-deploy-69cffd55cc-g999w       1/1     Running   0           17s
pod/test-deploy-69cffd55cc-gjjrk       1/1     Running   0           17s
pod/test-deploy-69cffd55cc-rlj67       1/1     Running   0           17s

NAME                                     TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes                      ClusterIP      10.96.0.1    <none>         443/TCP    76m

NAME                                     READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/test-deploy             3/3     3            0           18s

NAME                                     DESIRED   CURRENT   READY   AGE
replicaset.apps/test-deploy-69cffd55cc  3         3         3       18s
```

```
[root@master ~]# kubectl describe rs
Name:          test-deploy-69cffd55cc
Namespace:     default
Selector:      environment=test,pod-template-hash=69cffd55cc
Labels:        environment=test
               pod-template-hash=69cffd55cc
Annotations:   deployment.kubernetes.io/desired-replicas: 3
               deployment.kubernetes.io/max-replicas: 4
               deployment.kubernetes.io/revision: 1
Controlled By: Deployment/test-deploy
Replicas:      3 current / 3 desired
Pods Status:   3 Running / 0 Waiting / 0 Succeeded / 0 Failed
Pod Template:
  Labels:  environment=test
           pod-template-hash=69cffd55cc
  Containers:
    club-manager-application:
      Image:      azizur013/club-manager:latest
      Port:       80/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:      <none>
      Volumes:      <none>
Events:
  Type     Reason              Age   From                      Message
  ----     -
  Normal   SuccessfulCreate    7m26s replicaset-controller     Created pod: test-deploy-69cffd55cc-gjjrk
  Normal   SuccessfulCreate    7m26s replicaset-controller     Created pod: test-deploy-69cffd55cc-rlj67
  Normal   SuccessfulCreate    7m26s replicaset-controller     Created pod: test-deploy-69cffd55cc-g999w
```

## 7. Creating the Service:

#yml to create the services

apiVersion: v1

kind: Service

metadata:

name: test-deploy-service

spec:

selector:

environment: test

type: LoadBalancer

ports:

- protocol: TCP

port: 8000

targetPort: 80

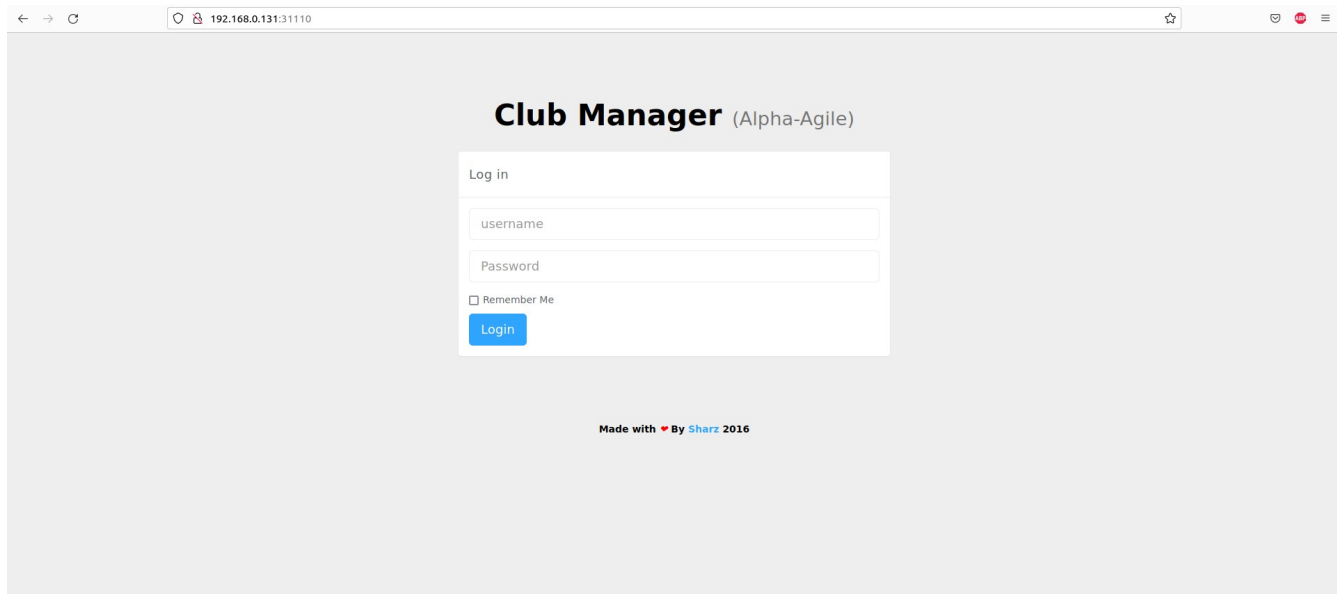
nodePort: 31110

```
[root@master ~]# kubectl apply -f club-manager-service.yml
service/test-deploy-service created
[root@master ~]# kubectl get sv
error: the server doesn't have a resource type "sv"
[root@master ~]# kubectl get svc
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)          AGE
kubernetes                         ClusterIP           10.96.0.1       <none>           443/TCP          100m
test-deploy-service                LoadBalancer       10.107.107.119  <pending>        8000:31110/TCP   67s
```

```
[root@master ~]# kubectl describe svc
Name:      kubernetet
Namespace: default
Labels:    component=apiserver
           provider=kubernetes
Annotations: <none>
Selector:   <none>
Type:       ClusterIP
IP Family Policy: SingleStack
IP Families: IPv4
IP:         10.96.0.1
IPs:        10.96.0.1
Port:       https 443/TCP
TargetPort: 6443/TCP
Endpoints:  192.168.0.131:6443
Session Affinity: None
Events:      <none>

Name:      test-deploy-service
Namespace: default
Labels:    <none>
Annotations: <none>
Selector:   environment=test
Type:       LoadBalancer
IP Family Policy: SingleStack
IP Families: IPv4
IP:         10.107.107.119
IPs:        10.107.107.119
Port:       <unset> 8000/TCP
TargetPort: 80/TCP
NodePort:   <unset> 31110/TCP
Endpoints:  10.10.1.2:80,10.10.1.3:80,10.10.1.4:80
Session Affinity: None
External Traffic Policy: Cluster
Events:      <none>
```

## 8.Final output :



## 9. References:

- a) Github Code Repo → <https://github.com/shindesharad71/Club-Manager.git>
- b) Docker image → <https://hub.docker.com/repository/docker/azizur013/club-manager>
- c) Kubeadm install documentation by Kubernetes → <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>
- d) Kubectl commands → <https://kubernetes.io/docs/reference/generated/kubectl/kubectl-commands>
- e) External code for centos 7 → <https://computingforgeeks.com/install-kubernetes-cluster-on-centos-with-kubeadm/>
- f) Kubernetes Dashboard configure → <https://computingforgeeks.com/how-to-install-kubernetes-dashboard-with-nodeport/>
- g) Youtube video reference → <https://www.youtube.com/watch?v=CfPRbdT-wXo>